

Strings Basics

STAT 133

Gaston Sanchez

`github.com/ucb-stat133/stat133-fall-2016`

Character Vectors Reminder

Character Basics

We express character strings using single or double quotes:

```
# string with single quotes  
'a character string using single quotes'
```

```
# string with double quotes  
"a character string using double quotes"
```

Character Basics

We can insert single quotes in a string with double quotes, and vice versa:

```
# single quotes within double quotes  
"The 'R' project for statistical computing"
```

```
# double quotes within single quotes  
'The "R" project for statistical computing'
```

Character Basics

We cannot insert single quotes in a string with single quotes, neither we can insert double quotes in a string with double quotes (Don't do this!):

```
# don't do this!  
"This "is" totally unacceptable"
```

```
# don't do this!  
'This 'is' absolutely wrong'
```

Function character()

Besides the single quotes or double quotes, R provides the function `character()` to create vectors of type character.

```
# character vector of 5 elements
```

```
a <- character(5)
```

```
a
```

```
## [1] "" "" "" "" ""
```

Empty string

The most basic string is the **empty string** produced by consecutive quotation marks: "".

```
# empty string  
empty_str <- ""  
  
empty_str  
  
## [1] ""
```

Technically, "" is a string with no characters in it, hence the name *empty string*.

Empty character vector

Another basic string structure is the **empty character vector** produced by `character(0)`:

```
# empty character vector  
empty_chr <- character(0)  
  
empty_chr  
  
## character(0)
```


Empty character vector

Do not to confuse the empty character vector `character(0)` with the empty string `""`; they have different lengths:

```
# length of empty string
```

```
length(empty_str)
```

```
## [1] 1
```

```
# length of empty character vector
```

```
length(empty_chr)
```

```
## [1] 0
```

Character Vectors

You can use the concatenate function `c()` to create character vectors:

```
strings <- c('one', '2', 'III', 'four')  
strings
```

```
## [1] "one"  "2"    "III"  "four"
```

```
example <- c('mon', 'tue', 'wed', 'thu', 'fri')  
example
```

```
## [1] "mon" "tue" "wed" "thu" "fri"
```

Replicate elements

You can also use the function `rep()` to create character vectors of replicated elements:

```
rep("a", times = 5)
rep(c("a", "b", "c"), times = 2)
rep(c("a", "b", "c"), times = c(3, 2, 1))
rep(c("a", "b", "c"), each = 2)
rep(c("a", "b", "c"), length.out = 5)
rep(c("a", "b", "c"), each = 2, times = 2)
```

Function paste()

The function `paste()` is perhaps one of the most important functions that we can use to create and build strings.

```
paste(..., sep = " ", collapse = NULL)
```

`paste()` takes one or more R objects, converts them to "character", and then it concatenates (pastes) them to form one or several character strings.

Function paste()

Simple example using paste():

```
# paste  
PI <- paste("The life of", pi)  
  
PI  
  
## [1] "The life of 3.14159265358979"
```

Function paste()

The default separator is a blank space (`sep = " "`). But you can select another character, for example `sep = "-"`:

```
# paste  
tobe <- paste("to", "be", "or", "not", "to", "be", sep = "-")  
  
tobe  
  
## [1] "to-be-or-not-to-be"
```

Function paste()

If we give paste() objects of different length, then the recycling rule is applied:

```
# paste with objects of different lengths
```

```
paste("X", 1:5, sep = ".")
```

```
## [1] "X.1" "X.2" "X.3" "X.4" "X.5"
```

Function paste()

To see the effect of the collapse argument, let's compare the difference with collapsing and without it:

```
# paste with collapsing  
paste(1:3, c("!", "?", "+"), sep = ' ', collapse = "")
```

```
## [1] "1!2?3+"
```

```
# paste without collapsing  
paste(1:3, c("!", "?", "+"), sep = ' ')
```

```
## [1] "1!" "2?" "3+"
```


Printing Strings

Printing Methods

Functions for printing strings can be very useful when creating our own functions. They help us have more control on the way the output gets printed either on screen or in a file.

Example str()

Many functions print output to the console. Some examples are `summary()` and `str()`:

```
# str
str(mtcars, vec.len = 1)

## 'data.frame': 32 obs. of 11 variables:
## $ mpg : num 21 21 ...
## $ cyl : num 6 6 ...
## $ disp: num 160 160 ...
## $ hp : num 110 110 ...
## $ drat: num 3.9 3.9 ...
## $ wt : num 2.62 ...
## $ qsec: num 16.5 ...
## $ vs : num 0 0 ...
## $ am : num 1 1 ...
## $ gear: num 4 4 ...
## $ carb: num 4 4 ...
```

Printing Characters

R provides a series of functions for printing strings.

Printing functions

Function	Description
<code>print()</code>	generic printing
<code>noquote()</code>	print with no quotes
<code>cat()</code>	concatenation
<code>format()</code>	special formats
<code>toString()</code>	convert to string
<code>sprintf()</code>	C-style printing

Method `print()`

The *workhorse* printing function in R is `print()`, which prints its argument on the console:

```
# text string
my_string <- "programming with data is fun"

# print string
print(my_string)

## [1] "programming with data is fun"
```

To be more precise, `print()` is a generic function, which means that you should use this function when creating printing methods for programmed classes.

Method `print()`

If we want to print character strings with no quotes we can set the argument `quote = FALSE`

```
# print without quotes  
print(my_string, quote = FALSE)  
  
## [1] programming with data is fun
```

Function noquote()

An alternative option for achieving a similar output is by using noquote()

```
# print without quotes
noquote(my_string)

## [1] programming with data is fun

# similar to:
print(my_string, quote = FALSE)

## [1] programming with data is fun
```

Function `cat()`

Another very useful function is `cat()` which allows us to concatenate objects and print them either on screen or to a file. Its usage has the following structure:

```
cat(..., file = "", sep = " ", fill = FALSE,  
     labels = NULL, append = FALSE)
```


Function `cat()`

If we use `cat()` with only one single string, you get a similar (although not identical) result as `noquote()`:

```
# simply print with 'cat()'  
cat(my_string)  
  
## programming with data is fun
```

`cat()` prints its arguments without quotes. In essence, `cat()` simply displays its content (on screen or in a file).

Function `cat()`

When we pass vectors to `cat()`, each of the elements are treated as though they were separate arguments:

```
# first four months  
cat(month.name[1:4], sep = " ")  
  
## January February March April
```

Function `cat()`

The argument `fill` allows us to break long strings; this is achieved when we specify the string width with an integer number:

```
# fill = 30
cat("Looooooooooooong strings", "can be displayed",
    "in a nice format",
    "by using the 'fill' argument", fill = 30)

## Looooooooooooong strings
## can be displayed
## in a nice format
## by using the 'fill' argument
```

Function `cat()`

Last but not least, we can specify a file output in `cat()`. For instance, to save the output in the file `output.txt` located in your working directory:

```
# cat with output in a given file  
cat(my_string, "with R", file = "output.txt")
```

Function `format()`

The function `format()` allows us to format an R object for pretty printing. This is especially useful when printing numbers and quantities under different formats.

```
# default usage
```

```
format(13.7)
```

```
## [1] "13.7"
```

```
# another example
```

```
format(13.12345678)
```

```
## [1] "13.12346"
```

Function `format()`

Some useful arguments of `format()`:

- ▶ `width` the (minimum) width of strings produced
- ▶ `trim` if set to `TRUE` there is no padding with spaces
- ▶ `justify` controls how padding takes place for strings.
Takes the values "left", "right", "centre", "none"

For controlling the printing of numbers, use these arguments:

- ▶ `digits` The number of digits to the right of the decimal place.
- ▶ `scientific` use `TRUE` for scientific notation, `FALSE` for standard notation

Function format()

```
# justify options
```

```
format(c("A", "BB", "CCC"), width = 5, justify = "centre")
```

```
## [1] "  A  " " BB  " " CCC "
```

```
format(c("A", "BB", "CCC"), width = 5, justify = "left")
```

```
## [1] "A      " "BB     " "CCC    "
```

```
format(c("A", "BB", "CCC"), width = 5, justify = "right")
```

```
## [1] "      A" "     BB" "    CCC"
```

```
format(c("A", "BB", "CCC"), width = 5, justify = "none")
```

```
## [1] "A"    "BB"   "CCC"
```

Function format()

```
# digits
format(1/1:5, digits = 2)

## [1] "1.00" "0.50" "0.33" "0.25" "0.20"

# use of 'digits', widths and justify
format(format(1/1:5, digits = 2), width = 6, justify = "c")

## [1] " 1.00" " 0.50" " 0.33" " 0.25" " 0.20"
```


string formatting with `sprintf()`

The function `sprintf()` is a wrapper for the C function `sprintf()` that returns a formatted string combining text and variable values. Its usage has the following form:

```
sprintf(fmt, ...)
```

The nice feature about `sprintf()` is that it provides us a very flexible way of formatting vector elements as character strings.

Using sprintf()

Several ways in which the number pi can be formatted:

```
# "%f" indicates 'fixed point' decimal notation
```

```
sprintf("%f", pi)
```

```
## [1] "3.141593"
```

```
# decimal notation with 3 decimal digits
```

```
sprintf("%.3f", pi)
```

```
## [1] "3.142"
```

```
# 1 integer and 0 decimal digits
```

```
sprintf("%1.0f", pi)
```

```
## [1] "3"
```

Using sprintf()

Several ways in which the number pi can be formatted:

```
# more options  
sprintf("%5.1f", pi)  
  
## [1] " 3.1"  
  
sprintf("%05.1f", pi)  
  
## [1] "003.1"
```

Using sprintf()

```
# print with sign (positive)
```

```
sprintf("%+f", pi)
```

```
## [1] "+3.141593"
```

```
# prefix a space
```

```
sprintf("% f", pi)
```

```
## [1] " 3.141593"
```

```
# left adjustment
```

```
sprintf("%-10f", pi) # left justified
```

```
## [1] "3.141593  "
```

Using sprintf()

```
# exponential decimal notation "e"  
sprintf("%e", pi)
```

```
## [1] "3.141593e+00"
```

```
# exponential decimal notation "E"  
sprintf("%E", pi)
```

```
## [1] "3.141593E+00"
```

```
# number of significant digits (6 by default)  
sprintf("%g", pi)
```

```
## [1] "3.14159"
```

Using sprintf()

```
# more sprintf examples  
sprintf("Harry's age is %s", 12)  
  
## [1] "Harry's age is 12"  
  
sprintf("five is %s, six is %s", 5, 6)  
  
## [1] "five is 5, six is 6"
```

Comparing printing methods

```
# printing method  
print(1:5)  
# convert to character  
as.character(1:5)  
# concatenation  
cat(1:5, sep="-")  
# default pasting  
paste(1:5)  
# paste with collapsing  
paste(1:5, collapse = "")  
# convert to a single string  
toString(1:5)  
# unquoted output  
noquote(as.character(1:5))
```

ggplot2 object

```
library(ggplot2)

gg <- ggplot(data = mtcars, aes(x = mpg, y = hp)) +
  geom_point()

summary(gg)

## data: mpg, cyl, disp, hp, drat, wt, qsec, vs, am, gear, carb [32x11]
## mapping: x = mpg, y = hp
## faceting: facet_null()
## -----
## geom_point: na.rm = FALSE
## stat_identity: na.rm = FALSE
## position_identity
```


ggplot2 summary()

<https://github.com/hadley/ggplot2/blob/master/R/summary.r>

```
summary.ggplot <- function(object, ...) {  
  wrap <- function(x) paste(  
    paste(strwrap(x, exdent = 2), collapse = "\n"), "\n", sep = "")  
  
  if (!is.null(object$data)) {  
    output <- paste(  
      "data:      ", paste(names(object$data), collapse = ", "),  
      " [", nrow(object$data), "x", ncol(object$data), "] ",  
      "\n", sep = "")  
    cat(wrap(output))  
  }  
  if (length(object$mapping) > 0) {  
    cat("mapping:  ", clist(object$mapping), "\n", sep = "")  
  }  
  if (object$scales$n() > 0) {  
    cat("scales:   ", paste(object$scales$input(), collapse = ", "), "\n")  
  }  
  cat("faceting: ")  
  print(object$facet)  
  if (length(object$layers) > 0)  
    cat("-----\n")  
  invisible(lapply(object$layers, function(x) {  
    print(x)  
    cat("\n")  
  })))  
}
```

Reading Raw Text

Reading Text with `readlines()`

- ▶ `readLines()` allows us to import text *as is* (i.e. we want to read raw text)
- ▶ Use `readLines()` if you don't want R to assume that the data is any particular form
- ▶ `readLines()` takes the name of a file or the name of a URL that we want to read
- ▶ The output is a character vector with one element for each line of the file or url

Reading Text with readlines()

For instance, here's how to read the file located at:

<http://www.textfiles.com/music/ktop100.txt>

```
# read 'ktop100.txt' file
ktop <- "http://www.textfiles.com/music/ktop100.txt"

top105 <- readLines(ktop)
```

Reading Text with readlines()

```
head(top105, n = 5)
```

```
## [1] "From: ed@wente.llnl.gov (Ed Suranyi)"  
## [2] "Date: 12 Jan 92 21:23:55 GMT"  
## [3] "Newsgroups: rec.music.misc"  
## [4] "Subject: KITS' year end countdown"  
## [5] ""
```

Basic String Manipulation

String Manipulation

There are a number of very handy functions in R for doing some basic manipulation of strings:

Manipulation of strings

Function	Description
<code>nchar()</code>	number of characters
<code>tolower()</code>	convert to lower case
<code>toupper()</code>	convert to upper case
<code>casefold()</code>	case folding
<code>chartr()</code>	character translation
<code>abbreviate()</code>	abbreviation
<code>substring()</code>	substrings of a character vector
<code>substr()</code>	substrings of a character vector

Counting number of characters

`nchar()` counts the number of characters in a string, that is, the “length” of a string:

```
# how many characters?  
nchar(c("How", "many", "characters?"))
```

```
## [1] 3 4 11
```

```
# how many characters?  
nchar("How many characters?")
```

```
## [1] 20
```

Notice that the white spaces between words in the second example are also counted as characters.

Counting number of characters

Do not confuse `nchar()` with `length()`. The former gives us the **number of characters**, the later only gives the **number of elements** in a vector:

```
# how many elements?  
length(c("How", "many", "characters?"))
```

```
## [1] 3
```

```
# how many elements?  
length("How many characters?")
```

```
## [1] 1
```

Convert to lower case with `tolower()`

R comes with three functions for text casefolding. The first function we'll discuss is `tolower()` which converts any upper case characters into lower case:

```
# to lower case  
tolower(c("aLL ChaRacterS in LowEr caSe", "ABCDE"))  
  
## [1] "all characters in lower case" "abcde"
```

Convert to upper case with toupper()

The opposite function of `tolower()` is `toupper`. As you may guess, this function converts any lower case characters into upper case:

```
# to upper case  
toupper(c("All ChaRacterS in Upper Case", "abcde"))  
  
## [1] "ALL CHARACTERS IN UPPER CASE" "ABCDE"
```

Case conversion with casefold()

casefold() converts all characters to lower case, but we can use the argument upper = TRUE to indicate the opposite (characters in upper case):

```
# lower case folding
casefold("aLL ChaRacterS in LoweR caSe")

## [1] "all characters in lower case"

# upper case folding
casefold("All ChaRacterS in Upper Case", upper = TRUE)

## [1] "ALL CHARACTERS IN UPPER CASE"
```

Character translation with `chartr()`

There's also the function `chartr()` which stands for *character translation*.

```
# character translation  
chartr(old, new, x)
```

`chartr()` takes three arguments: an old string, a new string, and a character vector `x`

Character translation with `chartr()`

The way `chartr()` works is by replacing the characters in `old` that appear in `x` by those indicated in `new`. For example, suppose we want to translate the letter "a" (lower case) with "A" (upper case) in the sentence `x`:

```
# replace 'a' by 'A'  
chartr("a", "A", "This is a boring string")  
  
## [1] "This is A boring string"
```

Character translation with `chartr()`

```
# multiple replacements
crazy <- c("Here's to the crazy ones",
           "The misfits",
           "The rebels")

chartr("aei", "#!?", crazy)

## [1] "H!r!'s to th! cr#zy on!s" "Th! m?sf?ts"
## [3] "Th! r!b!ls"
```

Abbreviate strings with `abbreviate()`

Another useful function for basic manipulation of character strings is `abbreviate()`. Its usage has the following structure:

```
abbreviate(names.org, minlength = 4, dot = FALSE,  
           strict =FALSE,  
           method = c("left.keep", "both.sides"))
```

Although there are several arguments, the main parameter is the character vector (`names.org`) which will contain the names that we want to abbreviate

Abbreviate strings with `abbreviate()`

```
# some color names
some_colors <- colors()[1:4]

# abbreviate (default usage)
colors1 <- abbreviate(some_colors)
colors1
```

##	white	aliceblue	antiquewhite	antiquewhite1
##	"whit"	"alcb"	"antq"	"ant1"

Abbreviate strings with `abbreviate()`

```
# abbreviate with 'minlength'
```

```
colors2 <- abbreviate(some_colors, minlength = 5)
```

```
colors2
```

```
##           white      aliceblue  antiquewhite  antiquewhite1
```

```
##      "white"      "alcbl"      "antqw"      "antq1"
```

```
# abbreviate
```

```
colors3 <- abbreviate(some_colors, minlength = 3,  
                      method = "both.sides")
```

```
colors3
```

```
##           white      aliceblue  antiquewhite  antiquewhite1
```

```
##      "wht"      "alc"      "ant"      "an1"
```

Replace substrings with substr()

The function `substr()` extracts or replaces substrings in a character vector. Its usage has the following form:

```
# replace  
substr(x, start, stop)
```

`x` is a character vector, `start` indicates the first element to be extracted (or replaced), and `stop` indicates the last element to be extracted (or replaced)

Extracting substrings with substr()

```
# extract characters in positions 1, 2, 3  
substr("abcdef", 1, 3)
```

```
## [1] "abc"
```

```
# extract 'area code'  
substr("(510) 987 6543", 2, 4)
```

```
## [1] "510"
```

Replace substrings with substr()

```
# replace 2nd letter with hash symbol
```

```
x <- c("may", "the", "force", "be", "with", "you")
```

```
substr(x, 2, 2) <- "#"
```

```
x
```

```
## [1] "m#y"    "t#e"    "f#rce"  "b#"     "w#th"   "y#u"
```

Replace substrings with substr()

```
# replace 2nd and 3rd letters with ":)"
y <- c("may", "the", "force", "be", "with", "you")

substr(y, 2, 3) <- ":)"

y

## [1] "m:)" "t:)" "f:)ce" "b:" "w:)h" "y:)"
```

Replace substrings with substr()

```
# replacement with recycling
```

```
z <- c("may", "the", "force", "be", "with", "you")
```

```
substr(z, 2, 3) <- c("#", "@")
```

```
z
```

```
## [1] "m#y"    "t@e"    "f#rce"  "b@"     "w#th"   "y@u"
```

Replace substrings with substring()

Closely related to `substr()`, the function `substring()` extracts or replaces substrings in a character vector. Its usage has the following form:

```
substring(text, first, last = 1000000L)
```

`text` is a character vector, `first` indicates the first element to be replaced, and `last` indicates the last element to be replaced

Replace substrings with substring()

```
# same as 'substr'  
substring("ABCDEF", 2, 4)  
  
## [1] "BCD"  
  
substr("ABCDEF", 2, 4)  
  
## [1] "BCD"  
  
# extract each letter  
substring("ABCDEF", 1:6, 1:6)  
  
## [1] "A" "B" "C" "D" "E" "F"
```

Replace substrings with substring()

```
# multiple replacement with recycling
txt <- c("another", "dummy", "example")

substring(txt, 1:3) <- c(" ", "zzz")

txt

## [1] " nother" "dzzzy"  "ex mple"
```

Set Operations

We can apply functions such as set union, intersection, difference, equality and membership, on "character" vectors.

Function	Description
<code>union()</code>	set union
<code>intersect()</code>	intersection
<code>setdiff()</code>	set difference
<code>setequal()</code>	equal sets
<code>identical()</code>	exact equality
<code>is.element()</code>	is element
<code>%in%()</code>	contains
<code>sort()</code>	sorting

Union

```
# two character vectors
set1 <- c("some", "random", "words", "some")

set2 <- c("some", "many", "none", "few")

# union of set1 and set2
union(set1, set2)

## [1] "some"    "random"  "words"   "many"    "none"    "few"
```

Intersection

```
# two character vectors
set3 <- c("some", "random", "few", "words")

set4 <- c("some", "many", "none", "few")

# intersect of set3 and set4
intersect(set3, set4)

## [1] "some" "few"
```

Set Difference

```
# two character vectors
set5 <- c("some", "random", "few", "words")

set6 <- c("some", "many", "none", "few")

# difference between set5 and set6
setdiff(set5, set6)

## [1] "random" "words"
```

Set Equality

```
# three character vectors
set7 <- c("some", "random", "strings")
set8 <- c("some", "many", "none", "few")
set9 <- c("strings", "random", "some")

# set7 == set8?
setequal(set7, set8)

## [1] FALSE

# set7 == set9?
setequal(set7, set9)

## [1] TRUE
```


Element Membership

```
# three vectors
set10 <- c("some", "stuff", "to", "play", "with")
elem1 <- "play"
elem2 <- "many"

# elem1 in set10?
is.element(elem1, set10)

## [1] TRUE

# elem2 in set10?
is.element(elem2, set10)

## [1] FALSE
```

Element Membership

```
# elem1 in set10?  
elem1 %in% set10
```

```
## [1] TRUE
```

```
# elem2 in set10?  
elem2 %in% set10
```

```
## [1] FALSE
```

Sorting

`sort()` arranges elements in alphabetical order

```
set11 <- c("random", "words", "multiple")
```

```
# sort (decreasingly)
```

```
sort(set11)
```

```
## [1] "multiple" "random"    "words"
```

```
# sort (increasingly)
```

```
sort(set11, decreasing = TRUE)
```

```
## [1] "words"    "random"   "multiple"
```

Package "stringr"

About "stringr"

About "stringr"

- ▶ functions are more consistent, simpler and easier to use
- ▶ "stringr" ensures that function and argument names (and positions) are consistent
- ▶ all functions deal with NA's and zero length character appropriately
- ▶ the output data structures from each function matches the input data structures of other functions

About "stringr"

"stringr" provides functions for both:

- ▶ basic manipulations and,
- ▶ for regular expression operations.

In this set of slides we cover those functions that have to do with basic manipulations.

About "stringr"

```
# installing 'stringr'  
install.packages("stringr")  
  
# load 'stringr'  
library(stringr)
```

Basic "stringr" functions

Function	Description	Similar to
<code>str_c()</code>	string concatenation	<code>paste()</code>
<code>str_length()</code>	number of characters	<code>nchar()</code>
<code>str_sub()</code>	extracts substrings	<code>substring()</code>
<code>str_dup()</code>	duplicates characters	<i>none</i>
<code>str_trim()</code>	removes leading and trailing whitespace	<i>none</i>
<code>str_pad()</code>	pads a string	<i>none</i>
<code>str_wrap()</code>	wraps a string paragraph	<code>strwrap()</code>
<code>str_trim()</code>	trims a string	<i>none</i>

About "stringr"

stringr provides functions for both:

- ▶ all functions in "stringr" start with `str_`
- ▶ some functions are designed to provide a better alternative to already existing functions
- ▶ Other functions don't have a corresponding alternative

Function str_c()

str_c() is equivalent to paste() but instead of using the white space as the default separator, str_c() uses the empty string ""

```
# default usage  
str_c("May", "The", "Force", "Be", "With", "You")  
  
## [1] "MayTheForceBeWithYou"
```

Function str_c()

Another major difference between str_c() and paste(): zero length arguments like NULL and character(0) are silently removed by str_c().

```
# removing zero length objects  
str_c("May", "The", "Force", NULL, "Be", "With", "You",  
      character(0))  
  
## [1] "MayTheForceBeWithYou"
```

Function str_c()

str_c() is equivalent to paste() but instead of using the white space as the default separator, str_c() uses the empty string ""

```
# changing separator
str_c("May", "The", "Force", "Be", "With", "You", sep="_")

## [1] "May_The_Force_Be_With_You"

# synonym function 'str_join'
str_join("May", "The", "Force", "Be", "With", "You", sep="-")

## Warning:  'str_join' is deprecated.
## Use 'str_c' instead.
## See help("Deprecated")

## [1] "May-The-Force-Be-With-You"
```

Function str_length()

str_length() is equivalent to nchar(), returning the number of characters in a string

```
# some text (NA included)
some_text = c("one", "two", "three", NA, "five")

# compare 'str_length' with 'nchar'
nchar(some_text)

## [1]  3  3  5 NA  4

str_length(some_text)

## [1]  3  3  5 NA  4
```

Function str_length()

str_length() has the nice feature that it converts factors to characters, something that nchar() is not able to handle:

```
# some factor
some_factor = factor(c(1, 1, 1, 2, 2, 2),
                     labels = c("good", "bad"))
some_factor

## [1] good good good bad  bad  bad
## Levels: good bad

# 'str_length' on a factor:
str_length(some_factor)

## [1] 4 4 4 3 3 3
```

Function str_length()

Compare str_length() against nchar()

```
# some factor
some_factor = factor(c(1,1,1,2,2,2),
                      labels = c("good", "bad"))

# now try 'nchar' on a factor
nchar(some_factor)

## Error in nchar(some_factor):  'nchar()' requires a
character vector
```

Function str_substr()

```
# some text
lorem = "Lorem Ipsum"

# apply 'str_sub'
str_sub(lorem, start=1, end=5)

## [1] "Lorem"

# equivalent to 'substring'
substring(lorem, first=1, last=5)

## [1] "Lorem"
```


Function str_substr()

str_sub() allows you to work with negative indices in the start and end positions:

```
# some strings
resto = c("brasserie", "bistrot", "creperie", "bouchon")

# 'str_sub' with negative positions
str_sub(resto, start=-4, end=-1)

## [1] "erie" "trot" "erie" "chon"
```

When we use a negative position, str_sub() counts backwards from last character.

Function str_sub()

A related function is `str_sub()`; when given a set of positions they will be recycled over the string

```
# extracting sequentially  
str_sub(lorem, seq_len(nchar(lorem)))  
  
## [1] "Lorem Ipsum" "orem Ipsum" "rem Ipsum" "em Ipsum"  
## [6] " Ipsum" "Ipsum" "psum" "sum"  
## [11] "m"
```

Function str_sub()

We can also give str_sub() a negative sequence, something that substring() ignores:

```
# reverse substrings with negative positions
```

```
str_sub(lorem, -seq_len(nchar(lorem)))
```

```
## [1] "m"          "um"          "sum"          "psum"          "  
## [6] " Ipsum"      "m Ipsum"      "em Ipsum"      "rem Ipsum"      "  
## [11] "Lorem Ipsum"
```

Function str_sub()

We can use str_sub() not only for extracting substrings but also for replacing substrings:

```
# replacing 'Lorem' with 'Nullam'  
lorem <- "Lorem Ipsum"  
str_sub(lorem, 1, 5) <- "Nullam"  
lorem  
  
## [1] "Nullam Ipsum"
```

Function str_sub()

```
# replacing with negative positions
lorem = "Lorem Ipsum"
str_sub(lorem, -1) <- "Nullam"
lorem

## [1] "Lorem IpsuNullam"

# multiple replacements
lorem = "Lorem Ipsum"
str_sub(lorem, c(1,7), c(5,8)) <- c("Nullam", "Enim")
lorem

## [1] "Nullam Ipsum" "Lorem Enimsum"
```

Duplication with str_dup()

str_dup() duplicates and concatenates strings within a character vector:

```
# default usage
str_dup("hola", 3)

## [1] "holaholahola"

# use with different 'times'
str_dup("adios", 1:3)

## [1] "adios"          "adiosadios"     "adiosadiosadios"
```

Duplication with str_dup()

```
# use with a string vector
words <- c("lorem", "ipsum", "dolor")
str_dup(words, 2)

## [1] "loremlorem" "ipsumipsum" "dolordolor"

str_dup(words, 1:3)

## [1] "lorem"          "ipsumipsum"     "dolordolordolor"
```

Padding with `str_pad()`

Another handy function that we can find in `stringr` is `str_pad()` for *padding* a string. Its default usage has the following form:

```
str_pad(string, width, side = "left", pad = " ")
```

The idea of `str_pad()` is to take a string and pad it with leading or trailing characters to a specified total width.

Padding with str_pad()

```
# default usage  
str_pad("hola", width=7)  
  
## [1] "    hola"  
  
# pad both sides  
str_pad("adios", width=7, side="both")  
  
## [1] " adios "
```

Padding with str_pad()

```
# left padding with '#'
```

```
str_pad("hashtag", width=8, pad="#")
```

```
## [1] "#hashtag"
```

```
# pad both sides with '-'
```

```
str_pad("hashtag", width=9, side="both", pad="-")
```

```
## [1] "-hashtag-"
```

Wrapping with `str_wrap()`

The function `str_wrap()` is equivalent to `strwrap()` which can be used to *wrap* a string to format paragraphs. Its default usage has the following form:

```
str_wrap(string, width = 80, indent = 0, exdent = 0)
```

Padding with `str_wrap()`

```
# quote (by Douglas Adams)
some_quote <- c(
  "I may not have gone",
  "where I intended to go,",
  "but I think I have ended up",
  "where I needed to be")

# some_quote in a single paragraph
some_quote <- paste(some_quote, collapse = " ")
```

Padding with `str_wrap()`

Say we want to display the text of `some_quote` within some pre-specified column width (e.g. width of 30):

```
# display paragraph with width=30  
cat(str_wrap(some_quote, width = 30))
```

```
## I may not have gone where I  
## intended to go, but I think I  
## have ended up where I needed  
## to be
```

Trimming with `str_trim()`

One of the typical tasks of string processing is that of parsing a text into individual words.

Usually, we end up with words that have blank spaces, called *whitespaces*, on either end of the word. In this situation, we can use the `str_trim()` function to remove any number of whitespaces at the ends of a string. Its usage requires only two arguments:

```
str_trim(string, side = "both")
```

Padding with `str_trim()`

```
# text with whitespaces
bad_text <- c(" several ", " whitespaces ")

# remove whitespaces on the left side
str_trim(bad_text, side = "left")

## [1] "several " "whitespaces "

# remove whitespaces on the right side
str_trim(bad_text, side = "right")

## [1] " several" " whitespaces"

# remove whitespaces on both sides
str_trim(bad_text, side = "both")

## [1] "several" "whitespaces"
```

Word extraction with word()

`word()` function that is designed to extract words from a sentence:

```
word(string, start = 1L, end = start, sep = fixed(" "))
```

The way in which we use `word()` is by passing it a string, together with a `start` position of the first word to extract, and an `end` position of the last word to extract. By default, the separator `sep` used between words is a single space.

Word extraction with word()

```
# some sentence
change = c("Be the change", "you want to be")

# extract first word
word(change, 1)

## [1] "Be"  "you"

# extract second word
word(change, 2)

## [1] "the"  "want"
```

Word extraction with word()

```
# some sentence
change = c("Be the change", "you want to be")

# extract last word
word(change, -1)

## [1] "change" "be"

# extract all but the first words
word(change, 2, -1)

## [1] "the change" "want to be"
```